

The goal is to implement a CNN to classify MNIST handwritten digit images using Python

Fetching the MNIST handwritten digit classification dataset-

```
# example of loading the mnist dataset
from tensorflow.keras.datasets import mnist
from matplotlib import pyplot as plt

# load dataset
(trainX, trainy), (testX, testy) = mnist.load_data()

# summarize loaded dataset
print('Train: X=%s, y=%s' % (trainX.shape, trainy.shape))
print('Test: X=%s, y=%s' % (testX.shape, testy.shape))

# plot first few images
for i in range(9):
    # define subplot
    plt.subplot(330 + 1 + i)

    # plot raw pixel data
    plt.imshow(trainX[i], cmap=plt.get_cmap('gray'))

# show the figure
plt.show()
```

Running the example loads the MNIST train and test dataset and prints their shape.

Although the MNIST dataset is effectively solved, it can be a useful starting point for developing and practicing a methodology for solving image classification tasks using convolutional neural networks.

We can load the images and reshape the data arrays to have a single-color channel-

```
# load dataset

(trainX, trainY), (testX, testY) = mnist.load_data()

# reshape dataset to have a single channel

trainX = trainX.reshape((trainX.shape[0], 28, 28, 1))

testX = testX.reshape((testX.shape[0], 28, 28, 1))
```

The load_dataset() function implements these behaviors and can be used to load the dataset-

```
# load train and test dataset

def load_dataset():

    # load dataset

    (trainX, trainY), (testX, testY) = mnist.load_data()

    # reshape dataset to have a single channel

    trainX = trainX.reshape((trainX.shape[0], 28, 28, 1))

    testX = testX.reshape((testX.shape[0], 28, 28, 1))

    # one hot encode target values

    trainY = to_categorical(trainY)

    testY = to_categorical(testY)

    return trainX, trainY, testX, testY
```

We know that the pixel values for each image in the dataset are unsigned integers in the range between black and white, or 0 and 255. We do not know the best way to scale the pixel values for modeling, but we know that some scaling will be required. Normalizing Pixel Values-

```
# convert from integers to floats

train_norm = train.astype('float32')

test_norm = test.astype('float32')

# normalize to range 0-1

train_norm = train_norm / 255.0

test_norm = test_norm / 255.0
```

Function must be called to prepare the pixel values prior to any modeling-

```
# scale pixels

def prep_pixels(train, test):

    # convert from integers to floats

    train_norm = train.astype('float32')

    test_norm = test.astype('float32')

    # normalize to range 0-1

    train_norm = train_norm / 255.0
```

```
test_norm = test_norm / 255.0

# return normalized images

return train_norm, test_norm
```

Next, we need to define a baseline convolutional neural network model for the problem-

```
# evaluate a model using k-fold cross-validation

def evaluate_model(dataX, dataY, n_folds=5):

    scores, histories = list(), list()

    # prepare cross validation

    kfold = KFold(n_folds, shuffle=True, random_state=1)

    # enumerate splits

    for train_ix, test_ix in kfold.split(dataX):

        # define model

        model = define_model()

        # select rows for train and test

        trainX, trainY, testX, testY = dataX[train_ix], dataY[train_ix], dataX[test_ix], dataY[test_ix]

        # fit model

        history = model.fit(trainX, trainY, epochs=10, batch_size=32, validation_data=(testX, testY), verbose=0)

        # evaluate model

        _, acc = model.evaluate(testX, testY, verbose=0)

        print('> %.3f' % (acc * 100.0))

        # stores scores

        scores.append(acc)

        histories.append(history)

    return scores, histories
```

Once the model has been evaluated, we can present the results-

```
# plot diagnostic learning curves

def summarize_diagnostics(histories):

    for i in range(len(histories)):

        # plot loss

        plt.subplot(2, 1, 1)

        plt.title('Cross Entropy Loss')

        plt.plot(histories[i].history['loss'], color='blue', label='train')

        plt.plot(histories[i].history['val_loss'], color='orange', label='test')

        # plot accuracy

        plt.subplot(2, 1, 2)
```

```
plt.title('Classification Accuracy')

plt.plot(histories[i].history['accuracy'], color='blue', label='train')

plt.plot(histories[i].history['val_accuracy'], color='orange', label='test')

plt.show()
```

The *summarize_performance()* function below implements this for a given list of scores collected during model evaluation-

```
# summarize model performance

def summarize_performance(scores):

# print summary

print('Accuracy: mean=%.3f std=%.3f, n=%d' % (mean(scores)*100, std(scores)*100, len(scores)))

# box and whisker plots of results

plt.boxplot(scores)

plt.show()
```

We need a function that will drive the test harness. This involves calling all the defined functions-

```
# run the test harness for evaluating a model

def run_test_harness():

# load dataset

trainX, trainY, testX, testY = load_dataset()

# prepare pixel data

trainX, testX = prep_pixels(trainX, testX)

# evaluate model

scores, histories = evaluate_model(trainX, trainY)

# learning curves

summarize_diagnostics(histories)

# summarize estimated performance

summarize_performance(scores)
```

We now have everything we need; the complete code example for a baseline convolutional neural network model on the MNIST dataset-

```
def load_dataset():

# load dataset

(trainX, trainY), (testX, testY) = mnist.load_data()

# reshape dataset to have a single channel

trainX = trainX.reshape((trainX.shape[0], 28, 28, 1))

testX = testX.reshape((testX.shape[0], 28, 28, 1))
```

```

# one hot encode target values
trainY = to_categorical(trainY)
testY = to_categorical(testY)
return trainX, trainY, testX, testY

# scale pixels
def prep_pixels(train, test):
    # convert from integers to floats
    train_norm = train.astype('float32')
    test_norm = test.astype('float32')

    # normalize to range 0-1
    train_norm = train_norm / 255.0
    test_norm = test_norm / 255.0

    # return normalized images
    return train_norm, test_norm

# define cnn model
def define_model():
    model = Sequential()
    model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform', input_shape=(28, 28, 1)))
    model.add(MaxPooling2D((2, 2)))
    model.add(Flatten())
    model.add(Dense(100, activation='relu', kernel_initializer='he_uniform'))
    model.add(Dense(10, activation='softmax'))

    # compile model
    opt = SGD(learning_rate=0.01, momentum=0.9)
    model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])

    return model

# evaluate a model using k-fold cross-validation
def evaluate_model(dataX, dataY, n_folds=5):
    scores, histories = list(), list()

    # prepare cross validation
    kfold = KFold(n_folds, shuffle=True, random_state=1)

    # enumerate splits
    for train_ix, test_ix in kfold.split(dataX):
        # define model
        model = define_model()

        # select rows for train and test

```

```

trainX, trainY, testX, testY = dataX[train_ix], dataY[train_ix], dataX[test_ix], dataY[test_ix]

# fit model

history = model.fit(trainX, trainY, epochs=10, batch_size=32, validation_data=(testX, testY), verbose=0)

# evaluate model

_, acc = model.evaluate(testX, testY, verbose=0)

print('> %.3f' % (acc * 100.0))

# stores scores

scores.append(acc)

histories.append(history)

return scores, histories

# plot diagnostic learning curves

def summarize_diagnostics(histories):

    for i in range(len(histories)):

        # plot loss

        plt.subplot(2, 1, 1)

        plt.title('Cross Entropy Loss')

        plt.plot(histories[i].history['loss'], color='blue', label='train')

        plt.plot(histories[i].history['val_loss'], color='orange', label='test')

        # plot accuracy

        plt.subplot(2, 1, 2)

        plt.title('Classification Accuracy')

        plt.plot(histories[i].history['accuracy'], color='blue', label='train')

        plt.plot(histories[i].history['val_accuracy'], color='orange', label='test')

        plt.show()

# summarize model performance

def summarize_performance(scores):

    # print summary

    print('Accuracy: mean=%.3f std=%.3f, n=%d' % (mean(scores)*100, std(scores)*100, len(scores)))

    # box and whisker plots of results

    plt.boxplot(scores)

    plt.show()

# run the test harness for evaluating a model

def run_test_harness():

    # load dataset

    trainX, trainY, testX, testY = load_dataset()

    # prepare pixel data

```

```
trainX, testX = prep_pixels(trainX, testX)

# evaluate model

scores, histories = evaluate_model(trainX, trainY)

# learning curves

summarize_diagnostics(histories)

# summarize estimated performance

summarize_performance(scores)

# entry point, run the test harness

run_test_harness()
```

DONE BY-

Sriman Satwik Reddy Chinnam

Yeruva Suprith Reddy

P. Rithvik Rao